

Reversible Session-Based Concurrency in Haskell*

Folkert de Vries¹ and Jorge A. Pérez¹[0000–0002–1452–6180]

University of Groningen, The Netherlands,

Abstract. Under a reversible semantics, computation steps can be undone. For message-passing, concurrent programs, reversing computation steps is a challenging and delicate task; one typically aims at formal semantics which are *causally-consistent*. Prior work has addressed this challenge in the context of a process model of multiparty protocols (choreographies) following a so-called *monitors-as-memories* approach. In this paper, we describe our ongoing efforts aimed at implementing this operational semantics in Haskell.

Keywords: Reversible computation · Message-passing concurrency · Session Types · Haskell.

1 Introduction

We implement the model in 1.

2 The Process Model

I think we need to explicitly define

- location
- participant
- queue

3 Our Haskell Implementation

We set out to implement the language, types and semantics given above. The end goal is to implement the two stepping functions

* F. de Vries is a BSc student.

```
forward :: Location -> Participant -> Session Value ()
backward :: Location -> Participant -> Session Value ()
```

Where `Session` contains an `ExecutionState` holding among other things a store of variables, and can fail producing an `Error`.

```
type Session value a = StateT (ExecutionState value) (Except Error) a
```

Additionally we need to provide a program for every participant, a monitor for every participant and a global message queue. All three of those need to be able to move forward and backward.

TODO list explicitly the next sections and what they describe

3.1 The Monitor

A participant is defined by its monitor and its program. The monitor contains various metadata about the participant: variables, the current state of the type and some other information to be able to move backward.

```
data Monitor value tipe =
  Monitor
    { _localType :: LocalTypeState (Program value) value tipe
    , _recursiveVariableNumber :: Int
    , _recursionPoints :: List (LocalType tipe)
    , _store :: Map Identifier value
    , _applicationHistory :: Map Identifier (Identifier, value)
    }
  deriving (Show, Eq)
```

Next we will look at how session types are represented, what the language looks like and how to keep track of and reverse past actions.

3.2 Global and Local Types

As mentioned, we have two kinds of session types: Global and Local. The Global type describes interactions between participants, specifically the sending and receiving of a value (a transaction), and selecting one out of a set of options (a choice). The definition of global types is given by

```
type GlobalType u = Fix (GlobalTypeF u)
```

```

data GlobalTypeF u next
  = Transaction
    { from :: Participant, to :: Participant, tipe :: u, continuation :: next }
  | Choice
    { from :: Participant, to :: Participant, options :: Map String next }
  | R next
  | V
  | Wk next
  | End
deriving (Show, Functor)

```

The recursive constructors are taken from 2. R introduces a recursion point, V jumps back to a recursion point and Wk weakens the recursion, making it possible to jump to a less tightly-binding R.

```

data MyParticipants = A | B | C | V deriving (Show, Eq, Ord)

data MyType = Title | Price | Share | Ok | Thunk | Address | Date
  deriving (Show, Eq, Ord)

globalType :: GlobalType.GlobalType MyParticipants MyType
globalType = GlobalType.globalType $ do
  GlobalType.transaction A V Title
  GlobalType.transactions V [A, B] Price
  GlobalType.transaction A B Share
  GlobalType.transactions B [A, V] Ok
  GlobalType.transaction B C Share
  GlobalType.transaction B C Thunk
  GlobalType.transaction B V Address
  GlobalType.transaction V B Date
  GlobalType.end

derivedTypeForA :: LocalType MyType
derivedTypeForA = do
  send V Title
  receive V Price
  send B Share
  receive B Ok

```

The Global type can then be projected onto a participant, resulting in a local type. The local type describes interactions between a participant and the central message queue. Specifically, sends and receives, and offers and selects. The projection of globalType onto A is equivalent to this pseudo-code of derivedTypeForA.

3.3 A Language

We need a language to use with our types. It needs at least instructions for the four participant-queue interactions, a way to assign variables, and a way to define and apply functions.

```
type Participant = String
type Identifier = String

type Program = Fix (ProgramF Value)

data ProgramF value next
  -- transaction primitives
  = Send { owner :: Participant, value :: value, continuation :: next }
  | Receive { owner :: Participant, variableName :: Identifier, continuation :: next }

  -- choice primitives
  | Offer Participant (List (String, next))
  | Select Participant (List (String, value, next))

  -- other constructors to make interesting examples
  | Parallel next next
  | Application Identifier value
  | Let Identifier value next
  | IfThenElse value next next
  | Literal value
  | NoOp
  deriving (Eq, Show, Functor)

data Value
  = VBool Bool
  | VInt Int
  | VString String
  | VUnit
  | VIntOperator Value IntOperator Value
  | VComparison Value Ordering Value
  | VFunction Identifier (Program Value)
  | VReference Identifier
  | VLabel String
  deriving (Eq, Show)
```

In the definition of `ProgramF`, the recursion is factored out and replaced by a type parameter. We then use `Fix` to give us back arbitrarily deep trees of instructions. The advantage of this transformation is that we can use recursion schemes - like folds - on the structure.

Given a `LocalType` and a `Program`, we can now step forward through the program. For each instruction, we check the session type to see whether the instruction is allowed.

3.4 An eDSL with the free monad

Writing programs with `Fix` everywhere is tedious, and we can do better. We can create an embedded domain-specific language (eDSL) using the free monad. The free monad is a monad that comes for free given some functor. With this monad we can use do-notation, which is much more pleasant to write.

The idea then is to use the free monad on our `ProgramF` data type to be able to build a nice DSL. For the transformation from `Free (ProgramF value)` a back to `Fix (ProgramF value)` we need also need some state: a variable counter that allows us to produce new unique variable names.

```
newtype HighLevelProgram a =
    HighLevelProgram (StateT (Location, Participant, Int) (Free (ProgramF Value))) a
    deriving (Functor, Applicative, Monad, MonadState (Location, Participant, Int))

uniqueVariableName :: HighLevelProgram Identifier
uniqueVariableName = do
    (location, participant, n) <- State.get
    State.put (location, participant, n + 1)
    return $ "var" ++ show n

send :: Value -> HighLevelProgram ()
send value = do
    (_, participant, _) <- State.get
    HighLevelProgram $ lift $ liftFree (Send participant value ())

receive :: HighLevelProgram Value
receive = do
    (_, participant, _) <- State.get
    variableName <- uniqueVariableName
    HighLevelProgram $ lift $ liftFree (Receive participant variableName ())
    return (VReference variableName)

terminate :: HighLevelProgram a
terminate = HighLevelProgram (lift $ Free NoOp)
```

We can now give correct implementations to the local types given above.

```
aType :: LocalType MyType
aType = do
```

```

    send V Title
    receive V Price
    send B Share
    receive B Ok

alice = H.compile "Location1" "A" $ do
  let share = VInt 42
  H.send (VString "address" )
  price <- H.receive
  H.send share
  ok <- H.receive
  H.terminate

```

And then transform them into a Program with

```

freeToFix :: Free (ProgramF value) a -> Program
freeToFix (Pure n) = Fix NoOp
freeToFix (Free x) = Fix (fmap freeToFix x)

compile :: Location -> Participant -> HighLevelProgram a -> Program
compile location participant (HighLevelProgram program) =
  freeToFix $ runStateT program (location, participant, 0)

```

3.5 Ownership

The owner field for send, receive, offer and select is important. It makes sure that instructions in closures are attributed to the correct participant.

```

bob = H.compile "Location1" "B" $ do
  thunk <-
    H.function $ \_ -> do
      H.send (VString "Lucca, 55100")
      d <- H.receive
      H.terminate

  price <- H.receive
  share <- H.receive
  let verdict = price `H.lessThan` VInt 79
  H.send verdict
  H.send verdict
  H.send share
  H.send thunk

```

```

carol = H.compile "Location1" "C" $ do
  h <- H.receive
  code <- H.receive
  H.applyFunction code VUnit

```

Here B creates a function that performs a send and receive. Because the function is created by B, the owner of these statements is B, even when the function is sent to and eventually evaluated by C.

The design of the language and semantics poses some further issues. With the current mechanism of storing applications, functions have to be named. Hence `H.function` cannot produce a simple value, because it needs to assign to a variable and thereby update the state.

It is also very important that all references in the function body are dereferenced before sending. Otherwise the function could fail on the other end or use variables that should be out of its scope.

3.6 Reversibility

Every forward step needs an inverse. When taking a forward step we store enough information to recreate the instruction and local type that made us perform the forward step.

```

type TypeContext program value a = Fix (TypeContextF program value a)

data TypeContextF program value a f
= Hole
| SendOrReceive (LocalTypeF a ()) f
| Selected
  { owner :: Participant
  , offerer :: Participant
  , selection :: Zipper (String, value, program, LocalType a)
  , continuation :: f
  }
| Offered
  { owner :: Participant
  , selector :: Participant
  , picked :: Zipper (String, program, LocalType a)
  , continuation :: f
  }
| Branched
  { condition :: value
  , verdict :: Bool
  , otherBranch :: program
  }

```

```

        , continuation :: f
      }
    | Application Identifier Identifier f
    | Assignment
      { visibleName :: Identifier
      , internalName :: Identifier
      , continuation :: f
      }
    | Literal a f
  deriving (Eq, Show, Generic, Functor)

```

For the instructions that modify the queue we must also roll the queue. Additionally, we require the their participants to be synchronized. Synchronization ensures that the complete transaction in the global type is undone, but the rolling can still happen in a decoupled way. The synchronization is a dynamic check that will give an error message if either participant is not in the expected state.

Let bindings remove assigned variables from the store. This is not strictly necessary to maintain causal consistency but it is good practice.

Function applications are treated exactly as in the formal semantics: We store a reference to the function and its arguments, so we can recreate the application later.

Given a `LocalType` and a `Program` we can now move forward whilst producing a trace through the execution. At any point, we can move back to a previous state.

3.7 Putting it all together

At the start we described the `Session` type.

```
type Session value a = StateT (ExecutionState value) (Except Error) a
```

We now have all the pieces we need to define the execution state. Based on the error conditions that arise in moving forward and backward, we can also define a meaningful `Error` type.

3.7.1 ExecutionState The execution state contains the monitors and the programs at all locations. Additionally it contains a variable counter to generate unique new names, and the central message queue

```
type Queue a = [a]
```



```

data ExecutionState value =
  ExecutionState
    { _variableCount :: Int
    , _applicationCount :: Int
    , _participants :: Map Participant (Monitor value String)
    , _locations :: Map Location (Map Participant (Program value))
    , _queue :: Queue value
    , _isFunction :: value -> Maybe (Identifier, Program value)
    }

```

3.7.2 Error generation There are a lot of potential failure conditions in this system. A small error somewhere in either the global type or the program can quickly move program and type out of sync. Therefore, returning detailed error messages is required.

```

data Error
  = SessionNotInSync
  | UndefinedParticipant Participant
  | UndefinedVariable Participant Identifier
  | SynchronizationError String
  | LabelError String
  | QueueError String Queue.QueueError
deriving (Eq, Show)

```

3.7.3 ?

```

forward :: Location -> Participant -> Session ()
backward :: Location -> Participant -> Session ()

```

4 Concluding Remarks and Future Work

- also store the current position in the global protocol and use it to step
- make informed decisions when a branch of a choice fails

Bibliography

- [1] C. A. Mezzina and J. A. Pérez, “Causally consistent reversible choreographies: A monitors-as-memories approach,” in *Proceedings of the 19th international symposium on principles and practice of declarative programming, namur, belgium, october 09 - 11, 2017*, 2017, pp. 127–138.
- [2] F. van Walree, “Session types in cloud haskell,” 2017.